

Analysis of Algorithms, I

CSOR W4231.002

Eleni Drinea
Computer Science Department

Columbia University

Tuesday, April 7, 2015

Outline

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow

Today

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow

Review of the last lecture


A union-find data structure for maintaining disjoint sets.

- ▶ Implementation: maintain sets as directed rooted trees;
 - ▶ **Makeset**: worst-case time is $O(1)$
 - ▶ **Find**: worst-case time is $O(\log n)$
 - ▶ **Union**: worst-case time is $O(\log n)$
- ▶ *Improved* implementation by using **path compression**; **amortized** time for a sequence of $2m$ **Find** operations: $O((m + n) \log^* n)$
 - ▶ this is not the tightest possible analysis but it is already fairly subtle

Today

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow

Modeling transportation networks



Source: Communications of the ACM, Vol. 57, No. 8

Can model a fluid network or a highway system by a graph:
edges carry *traffic*, nodes are *switches* where traffic gets diverted.

Flow networks

A flow network $G = (V, E)$ is a directed graph such that


1. Every edge has a capacity $c(e) \geq 0$. *A1: integer capacities*
2. There is a single source $s \in V$. *A2: no edge enters s*
3. There is a single sink $t \in V$. *A3: no edge leaves t*

Two more assumptions for the purposes of the analysis

- ▶ *A4: if $(u, v) \in E$ then $(v, u) \notin E$.*
- ▶ *A5: Every $v \in V - \{s, t\}$ is on some s - t path.*

Hence G has $m \geq n - 1$ edges.

An example flow network



Given a flow network G , an s - t flow f in G is a function

$$f : E \rightarrow R^+$$

Intuitively, the flow $f(e)$ on edge e is the amount of *traffic* that edge e carries.

Two kinds of constraints that every flow must satisfy

1. **Capacity constraints:** for all $e \in E$, $0 \leq f(e) \leq c(e)$.
2. **Flow conservation:** for all $v \in V - \{s, t\}$,

$$\sum_{(u,v) \in E} f(u,v) = \sum_{(v,w) \in E} f(v,w) \quad (1)$$

In words, the flow **into** node v equals the flow **out of** v , or

$$\sum_{e \text{ into } v} f(e) = \sum_{e \text{ out of } v} f(e)$$

A cleaner equation for flow conservation constraints

Define

$$1. f^{\text{out}}(v) = \sum_{e \text{ out of } v} f(e)$$

$$2. f^{\text{in}}(v) = \sum_{e \text{ into } v} f(e)$$

So we can rewrite equation (1) as: for all $v \in V - \{s, t\}$

$$f^{\text{in}}(v) = f^{\text{out}}(v) \tag{2}$$

The value of a flow


Definition 1.

The **value** of a flow f , denoted by $|f|$, is

$$|f| = \sum_{e \text{ out of } s} f(e) = f^{\text{out}}(s)$$


Can show that $|f| = f^{\text{in}}(t)$. (*exercise*)

An example flow of value 20



A flow f of value 20.

Another flow, of value 30



A (max) flow of value 30.

Max flow problem


Input: (G, s, t, c) such that

- ▶ $G = (V, E)$ is a flow network;
- ▶ $s, t \in V$ are the source and sink respectively;
- ▶ c is the (integer-valued) capacity function.

Output: a flow of maximum possible value


Definition 2.

An s - t cut (S, T) in G is a partition of the vertices into two sets S and T , such that $s \in S$ and $t \in T$.



A natural upper bound for the max value of a flow

- ▶ Flow f **must cross** (S, T) to go from source s to sink t .
- ▶ So it uses some (at most all) of the capacity of the edges crossing this cut.



- ▶ So, intuitively, the value of the flow cannot exceed

$$\sum_{e \text{ out of } S} c(e)$$

Definition 3.

The capacity $c(S, T)$ of an s - t cut (S, T) is defined as

$$c(S, T) = \sum_{e \text{ out of } S} c(e).$$

△ Note asymmetry in the definition of $c(S, T)$!

So, *intuitively*, the value of the max flow is upper bounded by the capacity of *every* cut in the flow network, that is,

$$\max_f |f| \leq \min_{(S, T) \text{ cut in } G} c(S, T) \quad (3)$$


Applications of max-flow and min-cut

- ▶ Min-cut
 - ▶ find a set edges of smallest capacity whose deletion disconnects the network
- ▶ Max-flow
 - ▶ Bipartite matching (*next lecture*)
 - ▶ Airline scheduling
 - ▶ Baseball elimination
 - ▶ Distribution of goods to cities
 - ▶ Image segmentation
 - ▶ Survey design
 - ▶ ...

Today

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow



“Undoing” flow



A flow f of value 20.

Would like to **undo** 10 units of flow along (u, v) and divert it along (u, t) .

Pushing flow back



- ▶ Push back 10 units of flow along (v, u) .
- ▶ Send 10 more units from s to t along edges (s, v) , (v, u) , (u, t) .
- ▶ New flow f' (on the right) with value 30.

Pushing flow forward and backward

By pushing flow back on (v, u) we created an **s - t path** on which we are pushing flow

- ▶ **Forward**, on edges with leftover capacity (e.g, on (s, v))
- ▶ **Backward**, on edges that are already carrying flow so as to divert it to a different direction (e.g., on (u, v)).

The residual graph G_f



Definition 4.

Given flow network G and flow f , the residual graph G_f has

- ▶ the **same vertices** as G ;
- ▶ for every edge $e = (u, v) \in E$ such that $f(e) < c(e)$, an edge $e = (u, v)$ with capacity $c_f(e) = c(e) - f(e)$ (**forward** edge);
- ▶ for every edge $e = (u, v) \in E$ such that $f(e) > 0$, an edge $e^r = (v, u)$ with capacity $c_f(e^r) = f(e)$ (**backward** edge).

So G_f has $\leq 2m$ edges.

Example residual graph



Left: a flow f of value 20.

Right: the residual graph G_f for this flow.

Augmenting paths

The residual graph G_f provides a roadmap for augmenting f .

1. Let P be a simple s - t path in G_f .
2. Augment f by pushing extra flow on P .

*How much extra flow can we push on P **without violating capacity constraints** in G_f ?*

- Let $c(P)$ be the capacity of path P defined as the **minimum residual capacity** of **any** edge of P .

$$c(P) = \min_{e \in P} c_f(e)$$

- The maximum amount of flow we can safely push on **every** edge of P is $c(P)$.

The augmented flow f'

Let P be an augmenting path in the residual graph G_f .

Augmented flow f' is as follows:

1. For a **forward** edge $e \in P$

$$f'(e) = f(e) + c(P)$$

2. For a **backward** edge $e^r = (u, v) \in P$, let $e = (v, u) \in G$

$$f'(e) = f(e) - c(P)$$

3. For $e \in E$ but not in P , $f'(e) = f(e)$.

Fact 5 (1).

f' is a flow.

Pseudocode

```
Augment( $f, P$ )  
  for each edge  $(u, v) \in P$  do  
    if  $e = (u, v)$  is a forward edge then  
       $f'(e) = f(e) + c(P)$   
    else  
       $f'(v, u) = f(v, u) - c(P)$   
    end if  
  end for  
  Return  $f'$ 
```

Today

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow

The Ford-Fulkerson algorithm

```
Ford-Fulkerson(  $G = (V, E, c), s, t$  )  
  for all  $e \in E$  do  $f(e) = 0$   
  end for  
  while there is an  $s$ - $t$  path in  $G_f$  do  
    Let  $P$  be a simple  $s$ - $t$  path in  $G_f$   
     $f' = \text{Augment}(f, P)$   
    Update  $f = f'$   
    Update  $G_f = G_{f'}$   
  end while  
  Return  $f'$ 
```

Running time analysis

The algorithm **terminates** if the following claims are both true

1. **Claim 1:** every iteration of the while loop returns a flow increased by an integer amount; and
2. **Claim 2:** there is a finite upper bound to the flow.

Proof of Claim 2.

Let U be the largest edge capacity. Then

$$|f| \leq \sum_{e \text{ out of } s} c(e) \leq nU$$



f increases by an integer amount after $\text{Augment}(f, P)$

Proof of Claim 1.

It follows from the following facts.

Fact 6 (2).

During execution of the Ford-Fulkerson algorithm, the flow values $\{f(e)\}$ and the residual capacities in G_f are all integers.

Fact 7 (3).

Let f be a flow in G and P a simple s - t path in G_f with residual capacity $c(P) > 0$. Then after $\text{Augment}(f, P)$

$$|f'| = |f| + c(P) \geq |f| + 1.$$



f increases by an integer amount after $\text{Augment}(f, P)$

Proof of Fact 3.

Recall that $|f| = f^{\text{out}}(s)$.

1. Since P is an s - t path, it contains an edge out of s , say (s, u) .
2. Since P is simple, it does not contain any edge entering s (P is in G_f , where there are edges entering s !): otherwise, s would be visited again.
3. Since no edge enters s in G , (s, u) is a forward edge in G_f , thus the flow on this edge is updated to $f(s, u) + c(P) \geq f(s, u) + 1$.
4. Since no other edge going out of s is updated, it follows that the value of f' is $|f'| = |f| + c(P) \geq |f| + 1$.



Running time of Ford-Fulkerson


1. Fact 3 guarantees at most nU iterations.
2. The running time of each iteration is bounded as follows:
 - ▶ $O(m + n)$ to create G_f using adjacency list representation
 - ▶ $O(m + n)$ to run BFS or DFS to find the augmenting path
 - ▶ $O(n)$ for $\text{Augment}(f, P)$ since P has at most $n - 1$ edges \Rightarrow Hence one iteration requires $O(m)$ time

The running time of Ford-Fulkerson is $O(mnU)$.

Remark 1.

*This is a **pseudo-polynomial** time algorithm: it is **not** polynomial in the description of the input U*

Problems with pseudo-polynomial running times



Today

- 1 Recap
- 2 Flow networks
 - Applications
- 3 The residual graph and augmenting paths
- 4 The Ford-Fulkerson algorithm for max flow
- 5 Faster algorithms for max flow

- ▶ Can be made polynomial: use BFS instead of DFS
 - ▶ Edmonds-Karp: $O(nm^2)$
- ▶ Unit capacities: $O(\min(\sqrt{m}, n^{2/3})m)$
- ▶ Integral capacities: $O(\min(\sqrt{m}, n^{2/3})m \log(n^2/m) \log U)$
[GoldbergRao1998]
- ▶ Real capacities: $O(nm \log(n^2/m))$
 - ▶ **Improved:** $O(nm)$ [Orlin2013]